

Predicting Lung Cancer with Machine Learning

October 9, 2022

```
[ ]: import pandas as pd
      %config Completer.use_jedi = False
      import numpy as np
      import seaborn as sns
      import matplotlib.pyplot as plt
      from sklearn.linear_model import LogisticRegression
      from sklearn.metrics import accuracy_score, log_loss, hinge_loss
      from sklearn.model_selection import train_test_split, KFold
      from statsmodels.stats.outliers_influence import variance_inflation_factor
      from sklearn.svm import SVC
```

```
[ ]: # Read in the data stored in the file 'survey lung cancer.csv'
      df = pd.read_csv('survey lung cancer.csv')
      df.head(5)
```

```
[ ]:  GENDER  AGE  SMOKING  YELLOW_FINGERS  ANXIETY  PEER_PRESSURE  \
0      M    69      1           2           2           1
1      M    74      2           1           1           1
2      F    59      1           1           1           2
3      M    63      2           2           2           1
4      F    63      1           2           1           1

      CHRONIC DISEASE  FATIGUE  ALLERGY  WHEEZING  ALCOHOL CONSUMING  COUGHING  \
0                   1         2         1         2           2         2
1                   2         2         2         1           1         1
2                   1         2         1         2           1         2
3                   1         1         1         1           2         1
4                   1         1         1         2           1         2

      SHORTNESS OF BREATH  SWALLOWING DIFFICULTY  CHEST PAIN  LUNG_CANCER
0                        2                      2           2         YES
1                        2                      2           2         YES
2                        2                      1           2         NO
3                        1                      2           2         NO
4                        2                      1           1         NO
```

```
[ ]: # replace all M/F with 1/0
df.GENDER.replace(['M', 'F'], [1, 0], inplace=True)

# replace all YES/NO with 1/0
df.LUNG_CANCER.replace(['YES', 'NO'], [1, 0], inplace=True)
```

1 Data analysis

```
[ ]: print("Number of duplicates: ", df.duplicated().sum())
df = df.drop_duplicates()
print("Number of duplcates after drop ", df.duplicated().sum())
```

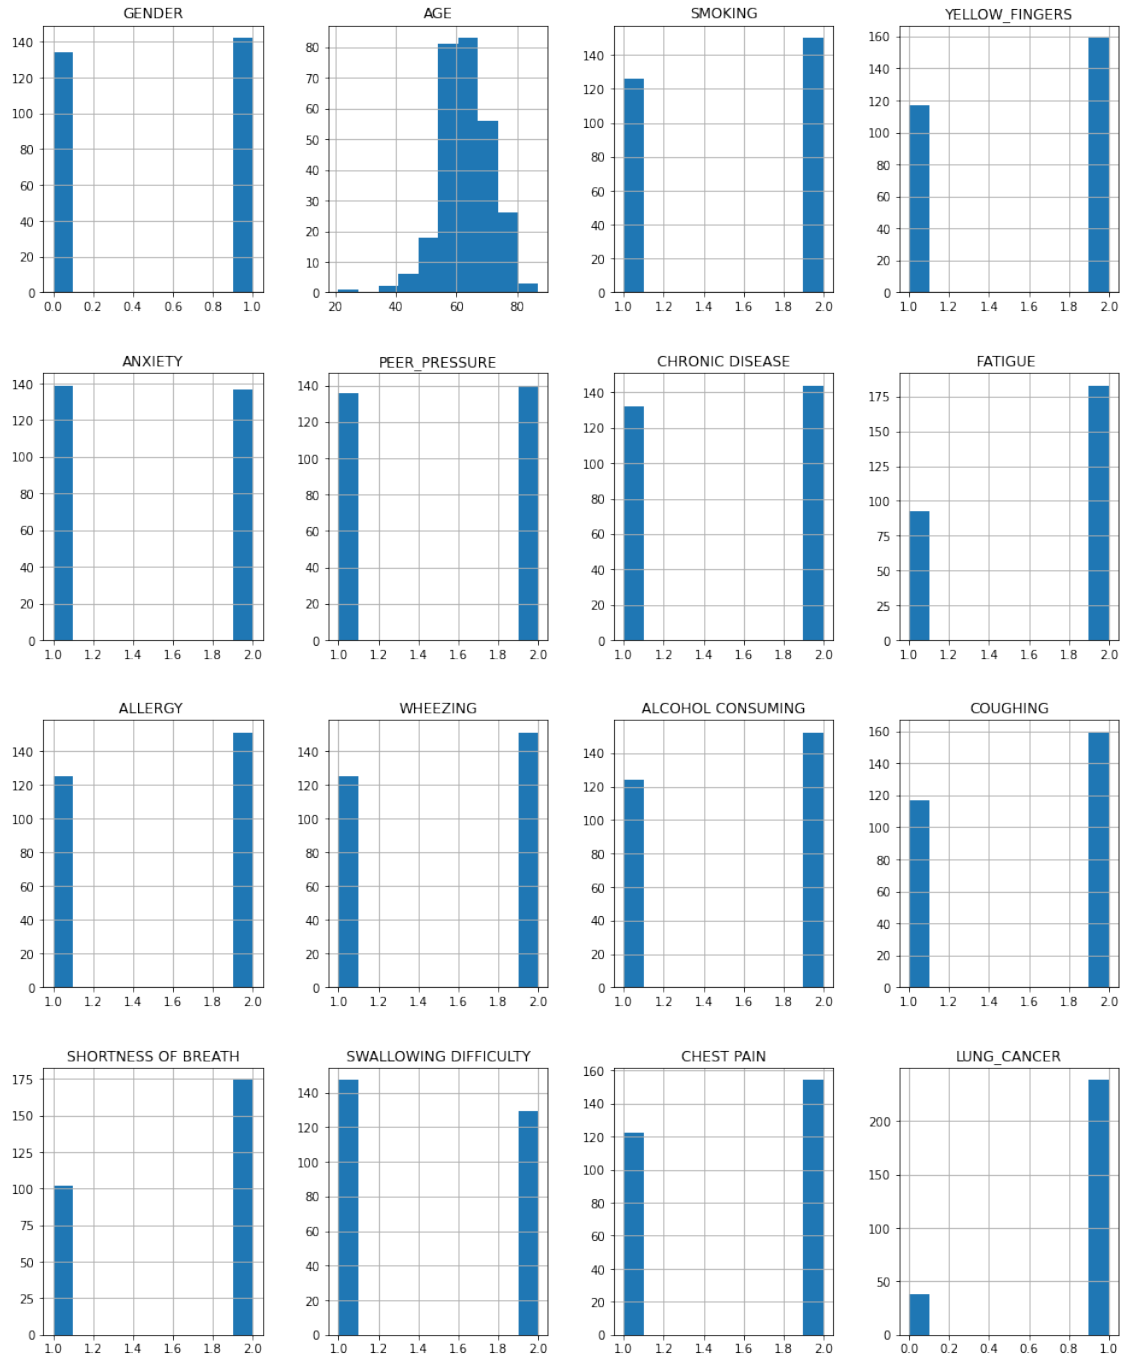
Number of duplicates: 33

Number of duplcates after drop 0

1.1 Histogram

```
[ ]: # Data histogram
df.hist(figsize=(16,20))
```

```
[ ]: array([[<AxesSubplot:title={'center':'GENDER'}>,
<AxesSubplot:title={'center':'AGE'}>,
<AxesSubplot:title={'center':'SMOKING'}>,
<AxesSubplot:title={'center':'YELLOW_FINGERS'}>],
[<AxesSubplot:title={'center':'ANXIETY'}>,
<AxesSubplot:title={'center':'PEER_PRESSURE'}>,
<AxesSubplot:title={'center':'CHRONIC DISEASE'}>,
<AxesSubplot:title={'center':'FATIGUE '}>],
[<AxesSubplot:title={'center':'ALLERGY '}>,
<AxesSubplot:title={'center':'WHEEZING'}>,
<AxesSubplot:title={'center':'ALCOHOL CONSUMING'}>,
<AxesSubplot:title={'center':'COUGHING'}>],
[<AxesSubplot:title={'center':'SHORTNESS OF BREATH'}>,
<AxesSubplot:title={'center':'SWALLOWING DIFFICULTY'}>,
<AxesSubplot:title={'center':'CHEST PAIN'}>,
<AxesSubplot:title={'center':'LUNG_CANCER'}>]], dtype=object)
```




1.2 Heatmap

```
[ ]: # create X and y sets
X = df.drop(columns="LUNG_CANCER").to_numpy()
X_ht = df.drop(columns="LUNG_CANCER", axis= 1)
y = df['LUNG_CANCER'].to_numpy().reshape(-1,)
```

```
[ ]: # Heatmap of dataframe
plt.figure(figsize=(16, 6))
mask = np.triu(np.ones_like(X_ht.corr()))
htmap = sns.heatmap(X_ht.corr(),vmin= -1, vmax= 1, annot= True , mask = mask)
htmap.set_title('Triangle Correlation Heatmap', fontdict={'fontsize':18},
    pad=16);

#https://medium.com/@szabo.bibor/
    how-to-create-a-seaborn-correlation-heatmap-in-python-834c0686b88e
```



1.3 Virtual Inflation Factor VIF

Was just used for some manual testing

```
[ ]: X_t = df.drop(columns="LUNG_CANCER", axis=1)
vif_data = pd.DataFrame()
vif_data["feature"] = X_t.columns
vif_data["VIF"] = [variance_inflation_factor(X_t.values, i) for i in
    range(len(X_t.columns))]
print(vif_data.sort_values(by = ["VIF"]))

#https://www.geeksforgeeks.org/detecting-multicollinearity-with-vif-python/
```

	feature	VIF
0	GENDER	3.012213
2	SMOKING	10.738372
6	CHRONIC DISEASE	10.896006
8	ALLERGY	13.490502

5	PEER_PRESSURE	13.564503
9	WHEEZING	13.986291
14	CHEST PAIN	14.003207
13	SWALLOWING DIFFICULTY	15.218347
11	COUGHING	17.344194
10	ALCOHOL CONSUMING	17.545363
12	SHORTNESS OF BREATH	17.666498
7	FATIGUE	17.770848
3	YELLOW_FINGERS	19.138132
4	ANXIETY	19.508213
1	AGE	41.961929

2 Machine Learning part

2.1 Data splitting

```
[ ]: # splitting out the test set, to be used later in comparing models
X_reduced, X_test, y_reduced, y_test = train_test_split(X, y, test_size=0.2,
↳shuffle=True)
```

2.2 Logistic Regression

```
[ ]: # Spitting the data in 5 sets
kf = KFold(n_splits=5, shuffle=True, random_state=69)

#results for train error
train_res_acc = []
train_res_err = []

# results for validation set
val_res_acc = []
val_res_err = []

# results for test set
test_res_acc = []
test_res_err = []

# iteration for k-fold
for train_index, val_index in kf.split(X_reduced):
    X_train, X_val= X_reduced[train_index], X_reduced[val_index]
    y_train, y_val= y_reduced[train_index], y_reduced[val_index]

    # make logistic regression model
    linreg = LogisticRegression(max_iter=100000)
    linreg.fit(X_train, y_train)
```

```

# predict training set
# calculate accuracy and logistic loss
y_pred_train = linreg.predict(X_train)
train_acc = accuracy_score(y_train, y_pred_train)
train_err = log_loss(y_train, y_pred_train)

# predict validation set
# calculate accuracy and logistic loss
y_pred_val = linreg.predict(X_val)
val_acc = accuracy_score(y_val, y_pred_val)
val_error = log_loss(y_val, y_pred_val)

# predict test set
# calculate accuracy and logistic loss
y_pred_test = linreg.predict(X_test)
test_acc = accuracy_score(y_test, y_pred_test)
test_error = log_loss(y_test, y_pred_test)

# append train results for later
train_res_acc.append(train_acc)
train_res_err.append(train_err)

# append validation results for later
val_res_acc.append(val_acc)
val_res_err.append(val_error)

# append test results for later
test_res_acc.append(test_acc)
test_res_err.append(test_error)

# print results
print("Logistic Regression")
print("Training accuracy: ", np.mean(train_res_acc))
print("Training error: ", np.mean(train_res_err))
print("Validation accuracy: ", np.mean(val_res_acc))
print("Validation error: ", np.mean(val_res_err))
print("Test accuracy: ", np.mean(test_res_acc))
print("Test error: ", np.mean(test_res_err))

```

```

Logistic Regression
Training accuracy:  0.9295454545454545
Training error:    2.433456497231457

```

Validation accuracy: 0.8954545454545455
Validation error: 3.610937499257678
Test accuracy: 0.9035714285714287
Test error: 3.3305734136389353

2.3 SVC

```
[ ]: # Spitting the data in 5 sets
kf = KFold(n_splits=5, shuffle=True, random_state=69)

#results for train error
train_res_acc2 = []
train_res_err2 = []

# results for validation set
val_res_acc2 = []
val_res_err2 = []

# results for test set
test_res_acc = []
test_res_err = []

# iteration for k-fold
for train_index, val_index in kf.split(X_reduced):
    X_train, X_val= X_reduced[train_index], X_reduced[val_index]
    y_train, y_val= y_reduced[train_index], y_reduced[val_index]

    #make SVC model
    svc = SVC(gamma='auto',kernel='linear')
    svc.fit(X_train,y_train)

    # predict training set
    # calculate accuracy and logistic loss
    y_pred_train2 = svc.predict(X_train)
    train_acc2 = accuracy_score(y_train, y_pred_train2)
    train_err2 = hinge_loss(y_train, y_pred_train2)

    # predict validation set
    # calculate accuracy and logistic loss
    y_pred_val2 = svc.predict(X_val)
    val_acc2 = accuracy_score(y_val, y_pred_val2)
    val_error2 = hinge_loss(y_val, y_pred_val2)

    # predict test set
    # calculate accuracy and logistic loss
    y_pred_test = svc.predict(X_test)
```

```

test_acc = accuracy_score(y_test, y_pred_test)
test_error = hinge_loss(y_test, y_pred_test)

# append train results for later
train_res_acc.append(train_acc2)
train_res_err.append(train_err2)

# append validation results for later
val_res_acc.append(val_acc2)
val_res_err.append(val_error2)

# append test results for later
test_res_acc.append(test_acc)
test_res_err.append(test_error)

# print results
print("SVC")
print("Training accuracy: ", np.mean(train_res_acc))
print("Training error: ", np.mean(train_res_err))
print("Validation accuracy: ", np.mean(val_res_acc))
print("Validation error: ", np.mean(val_res_err))
print("Test accuracy: ", np.mean(test_res_acc))
print("Test error: ", np.mean(test_res_err))

```

```

SVC
Training accuracy:  0.9363636363636363
Training error:    1.3155918849793649
Validation accuracy:  0.9
Validation error:    1.923650567810657
Test accuracy:      0.9214285714285715
Test error:         0.20357142857142857

```

[]: